

CSE221 Data Structures

Lecture 17: Ordered Maps and Skip Lists

Antoine Vigneron
`antoine@unist.ac.kr`

Ulsan National Institute of Science and Technology


November 15, 2021

- 1 Introduction
- 2 Ordered maps
- 3 Ordered search tables and binary search
- 4 Skip lists
- 5 Probabilistic analysis of skip lists

Introduction

- I updated attendance records. They can be found in the portal under E-attendance.
- I will post Assignment 3 tonight.
- Reference for this lecture: Textbook Chapter 9.3–9.4.

Ordered Maps



- An *ordered map* ADT is similar with a map ADT, but uses an ordering of the keys.
- In particular, it may use key *comparators*. (See Lecture 15.)

Ordered Maps

- In addition to the map ADT, the *ordered map* ADT supports the following functions:

- **firstEntry**(k): Return an iterator to the entry with smallest key value.
- **lastEntry**(k): Return an iterator to the entry with largest key value.
- **ceilingEntry**(k): Return an iterator to the entry with the least key value greater than or equal to k .
- **floorEntry**(k): Return an iterator to the entry with the greatest key value less than or equal to k .
- **lowerEntry**(k): Return an iterator to the entry with the greatest key value less than k .
- **higherEntry**(k): Return an iterator to the entry with the least key value greater than k .

(When there is no such entry, we return an iterator end.)

Ordered Maps

- Intuitively, a map implemented with a hash table would not be a good implementation of an ordered map, because the entries are not placed in any particular order in the table. In fact, the hash function is chosen in such a way that the positions look random.
- So we will need different implementation.

Example

- A database records flights in an ordered map, with keys

$k = (\text{origin}, \text{destination}, \text{date}, \text{time})$.

in lexicographical order.

- Suppose you want to take a flight from ICN to CDG after Nov 17, 9:30 am.
- You can find the 4 earliest such flights by calling `ceilingEntry(ICN, CDG, 17Nov, 09:30)`, followed by 3 calls to `higherEntry` on the successive results:





```
((ICN, CDG, 17Nov, 09:55), (AF237 ... ))  
((ICN, CDG, 17Nov, 15:15), (KE516 ... ))  
((ICN, CDG, 17Nov, 19:50), (OZ218 ... ))  
((ICN, CDG, 18Nov, 00:45), (AF354 ... ))
```

Ordered Search Tables

| | | | | | | | | | | |
|---|---|---|----|----|----|----|----|----|---|----|
| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 |
| 4 | 6 | 9 | 12 | 15 | 16 | 18 | 28 | 34 | | |

- An *ordered search table* is a vector L where the entries are stored by increasing key values.
- Assuming we grow and shrink L in an appropriate manner, the space usage is $O(n)$.
- Inserting an entry (k, v) takes $O(n)$ time as we may need to shift a linear number of entries.
- Deleting an element also takes $O(n)$ time.
- On the other hand, we will see that it allows us to perform other operations to run in $O(\log n)$ time.

Binary Search: Example with Find(22)



Binary Search

- We are given a vector L of n entries, with index 0 to $n - 1$.
- The `find` operation of the map ADT can be implemented by *binary search*.
- We recurse on an interval of indices $[\text{low}, \text{high}]$ where $0 \leq \text{low} \leq \text{high} \leq n - 1$. Initially, $\text{low} = 0$ and $\text{high} = n - 1$.
- We make sure that, if the key is in L , its index must be in the interval $[\text{low}, \text{high}]$.
- Let e be the entry with index

$$\text{mid} = \left\lfloor \frac{\text{low} + \text{high}}{2} \right\rfloor.$$

Binary Search

- If $k = e.\text{key}$, then we are done: We found k .
- If $k < e.\text{key}$, then we recurse on the range $[\text{low}, \text{mid}-1]$.
- If $k > e.\text{key}$, then we recurse on the range $[\text{mid}+1, \text{high}]$.

Pseudocode

```
procedure BINARYSEARCH( $L, k, \text{low}, \text{high}$ )  
  if  $\text{low} > \text{high}$  then  
    return end  
   $\text{mid} \leftarrow \lfloor (\text{low} + \text{high})/2 \rfloor$   
   $e \leftarrow L[\text{mid}]$   
  if  $k = e.\text{key}$  then  
    return  $e$   
  if  $k < e.\text{key}$  then  
    return BINARYSEARCH( $L, k, \text{low}, \text{mid}-1$ )  
  return BINARYSEARCH( $L, k, \text{mid}+1, \text{high}$ )
```

Analysis

- Each recursive call takes $O(1)$ time.
- Let $n_i = \text{high} - \text{low} + 1$ denote the size of the range under consideration at step i
- Initially, $n_0 = n$. At each step i , we have $n_{i+1} \leq n_i/2$.
- So at step i , we have $1 \leq n_i \leq n/2^i$.
- It follows that $2^i \leq n$ and thus $i \leq \log n$.
- In summary, there are at most $\log n$ steps, each step takes $O(1)$ time, so the running time of binary search is

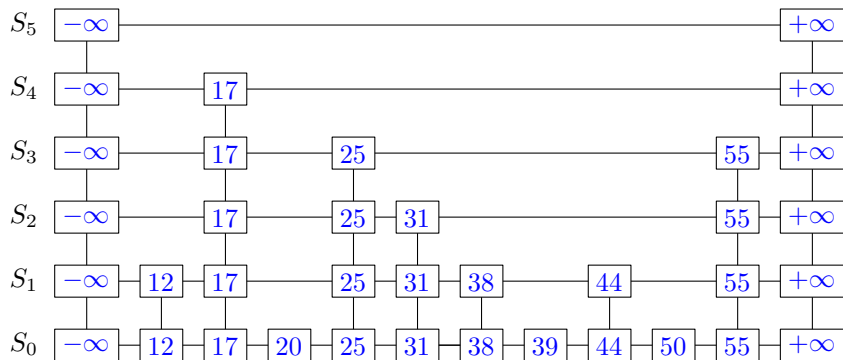
$$O(\log n).$$

Comparing Map Implementations

| method | List | Hash Table | Search Table |
|-------------|--------|--------------------------------|--------------|
| size, empty | $O(1)$ | $O(1)$ | $O(1)$ |
| find | $O(n)$ | $O(1)$ exp., $O(n)$ worst-case | $O(\log n)$ |
| insert | $O(1)$ | $O(1)$ | $O(n)$ |
| erase | $O(n)$ | $O(1)$ exp., $O(n)$ worst-case | $O(n)$ |

- We will now show that *skip lists* allow us to perform all the operations of the *ordered* map ADT efficiently.


Skip Lists




Skip Lists

- A *skip list* for a map M consists of a sequence of lists $S_0 \supseteq S_1 \supseteq \dots \supseteq S_h$.
- h is the *height* of the skip list.
- Each element of S_i appears in S_{i+1} with probability $1/2$.
- So the size of S_{i+1} is roughly half of the size of S_i .
- This is called *randomization*: We make random choices to construct the data structure.
- Then we will show that, *on average*, a search or update operation takes $O(\log n)$ time. This is an *expected* time bound, as opposed to worst-case time bounds for ordered search tables.

Skip Lists



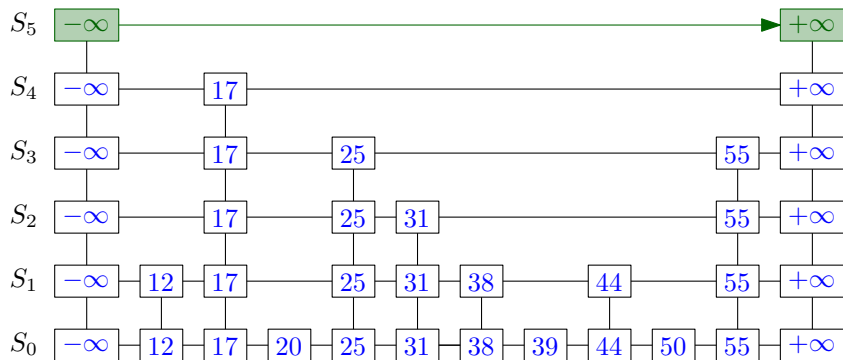
Skip Lists




Operations for traversing a skip list

- **after(p)**: Return the position following p on the same level.
 - **before(p)**: Return the position preceding p on the same level.
 - **below(p)**: Return the position below p in the same tower.
 - **above(p)**: Return the position above p in the same tower.
-
- Using a linked structure, each of these operations takes $O(1)$ time

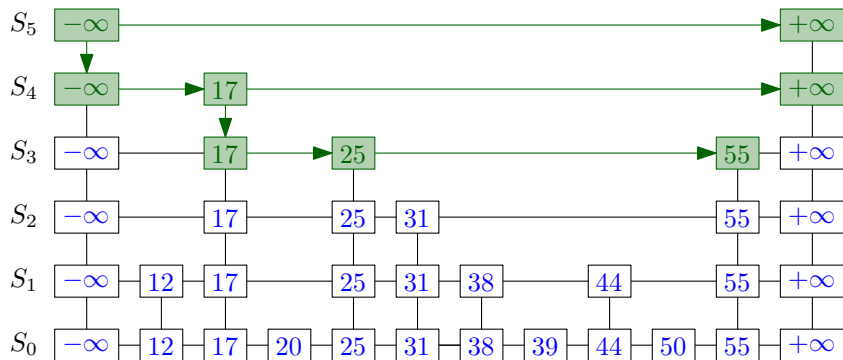
Searching for Key 50




Searching for Key 50




Searching for Key 50




Searching for Key 50



Searching for Key 50



Searching for Key 50




Searching in a Skip List

Pseudocode

```
1: procedure SKIPSEARCH( $k$ )
2:    $p \leftarrow s$ 
3:   while below( $p$ )  $\neq$  NULL do
4:      $p \leftarrow$  below( $p$ )                                ▷ drop down
5:     while  $k \geq$  key(after( $p$ )) do
6:        $p \leftarrow$  after( $p$ )                                ▷ scan forward
```


- Searching in a skip list S . Variable s holds the start position of S .
- This function returns the position p in the bottom list S_0 such that the entry at p has the largest key less than or equal to k .
- The operation $p \leftarrow$ below(p) is called *drop down*. We drop down until we reach the *bottom*, in which case we are done.
- At Line 5–6, we *scan forward*: We look for the right-most position p such that $\text{key}(p) \leq k$.

Inserting 42




- First perform $\text{SKIPSEARCH}(42)$.

Inserting 42




- Insert new node at the bottom.

Inserting 42




- With probability $1/2$: $\text{coinFlip()}=\text{tails}$. Grow the tower.

Inserting 42



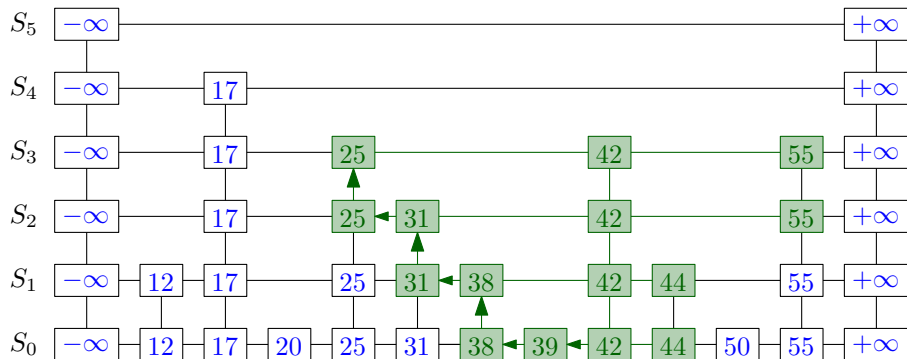
- $\text{coinFlip()}=\text{tails}$. Grow the tower.

Inserting 42



- $\text{coinFlip()}=\text{tails}$. Grow the tower.

Inserting 42




- coinFlip() \equiv heads: We are done


Skip Lists

- Detailed pseudocode is given in the textbook.
- Remark: If the tower grows beyond S_h , we need to add new levels to the skip list.


Removal of 25: First perform SKIPSEARCH(25).




Removal of 25




Removal of 25



Removal of 25



Removal of 25



Probabilistic Analysis of Skip Lists

- We first bound the height of a skip list.
- Let e be an entry of a skip list, and $h(e)$ the height of its tower.
- Then

$$\Pr[h(e) \geq i] = \Pr[\text{first } i \text{ coin flips are "tails"}] = 1/2^i$$

- The probability that level i is non-empty is

$$\begin{aligned}\Pr[S_i \neq \emptyset] &= \Pr[\text{there is an entry } e \text{ such that } h(e) \geq i] \\ &\leq \sum_{e \in S} \Pr[h(e) \geq i] \\ &= n \cdot (1/2^i).\end{aligned}$$

Probabilistic Analysis of Skip Lists

- So the probability that the height h of S is at most $3 \log n$ satisfies:

$$\begin{aligned}\Pr[h \geq 3 \log n] &\leq n \cdot \frac{1}{2^{3 \log n}} \\ &= n \cdot \frac{1}{n^3} = \frac{1}{n^2}.\end{aligned}$$

- We say that $h = O(\log n)$ with *high probability* because $h < c \log n$ with probability at least $1 - 1/n^{c-1}$, for some constant $c > 1$. (Here $c = 3$.)

Example

If S records 1,000 entries, then the probability that its height is more than $3 \log 1000 \approx 30$ is less than 10^{-6} .

Probabilistic Analysis of Skip Lists

- With high probability, $h = O(\log n)$.
- So in a search operation, the number of *drop-down* moves is $O(\log n)$.
- What about the number of *scan-forward*?
- Let n_i be the number of keys examined while scanning forward at level i .
- None of these n_i keys was in S_{i+1} , otherwise they would have been scanned at the previous step.
- Each key of S_i is in S_{i+1} with probability $1/2$.
- Therefore, the expected value of n_i is exactly equal to the expected number of times we must flip a fair coin before it comes up heads. This expected value is 2.

Probabilistic Analysis of Skip Lists

- So for each level, we perform 1 drop-down and an expected number of 2 scan-forward.
- Hence the expected running time is $O(h) = O(\log n)$ with high probability.
- It follows that the expected search time is $O(\log n)$.
- Similarly, we can prove that the expected insertion and removal times are $O(\log n)$.