

CSE221 Data Structures

Lecture 21: Directed Graphs

Antoine Vigneron
`antoine@unist.ac.kr`

Ulsan National Institute of Science and Technology

November 29, 2021

- 1 Introduction
- 2 Data structure
- 3 Reachability
- 4 Cycle detection
- 5 Strong connectivity
- 6 Directed acyclic graphs
- 7 Shortest Paths



Introduction

- Final exam is on Wednesday 15 December, 20:00–22:00.
- Assignment 4 will be posted tomorrow, due on Thursday next week.
- Assignment 3 will be graded by Seonghyeon Jue (shjue@unist.ac.kr)
- Today's lecture is on algorithms for *directed* graphs.
- Reference for this lecture: Textbook Chapter 13.4 and 13.5

Data Structure

- In addition to the standard ADT functions, we will need the following:
- **e.isDirected()**: Test whether edge e is directed.
- **e.origin()**: Return the origin vertex of edge e .
- **e.dest()**: Return the destination vertex of edge e .
- **insertDirectedEdge**(v, w, x): Insert and return a new directed edge with origin v and destination w and storing element x .
- It allows us to deal with *directed graphs*, or *mixed graphs* (that have undirected and directed edges).

Reachability



Definition

Given two vertices u , v in a graph G , we say that u *reaches* v if there is a directed path from u to v . We also say that v is *reachable* from u . The *reachability* problem is to decide whether u reaches v .

Example

In the graph above, u reaches v , but u is not reachable from v .

Reachability



- This problem can be solved by depth-first search.

Pseudocode

```
1: procedure DIRECTEDDFS( $v$ )  
2:   mark  $v$  as visited  
3:   for each outgoing edge  $(v, w)$  of  $v$  do  
4:     if vertex  $w$  has not been visited then  
5:       recursively call DIRECTEDDFS( $w$ )
```


- DFS visits *all* the vertices that are reachable from v .
- It is the same algorithms as for undirected graphs, except that we only move from the origin to the destination of an edge.
- It still runs in $O(n + m)$ time where n is the number of vertices and m is the number of edges. So we can solve the reachability problem in $O(n + m)$ time.

Reachability



- We can also run BFS on a directed graph in $O(n + m)$ time.
- It creates three types of edges: Tree edges, cross edges and back edges. Cross edges connect a vertex to a vertex that is neither its ancestor nor its descendent. Back edges that connect a vertex to one of its ancestors.
- In the example above, BFS produces one back edge (e_8) and two cross edges (e_6 and e_7).

Cycle Detection




Problem

Given a directed graph G , the **cycle detection** problem is to find a directed cycle in G , if there is at least one.

Example


In the graph above, $(v_4, e_6, v_5, e_7, v_6, e_8, v_4)$ is a cycle, so a cycle detection algorithm could just return this cycle. There is another cycle $(v_1, e_3, v_4, e_9, v_1)$, so the algorithm could have returned this one instead.

Cycle Detection




- For *undirected* graphs, we saw in the previous lecture that it suffices to run DFS, and if we encounter a previously visited vertex at Line 4, Slide 6, then we the edge under consideration closes a cycle.

Cycle Detection




- For *undirected* graphs, we saw in the previous lecture that it suffices to run DFS, and if we encounter a previously visited vertex at Line 4, Slide 6, then we the edge under consideration closes a cycle.

Cycle Detection




- For *undirected* graphs, we saw in the previous lecture that it suffices to run DFS, and if we encounter a previously visited vertex at Line 4, Slide 6, then we the edge under consideration closes a cycle.

Cycle Detection




- This approach does not work for *directed* graphs.

Cycle Detection




- This approach does not work for *directed* graphs.

Cycle Detection




- This approach does not work for *directed* graphs.

Cycle Detection



- This is not a directed cycle.

Cycle Detection




- A correct approach: Perform directed DFS.
- If a **back edge** is found, then it closes a cycle. (A back edge is an edge whose destination is an ancestor of its origin in the spanning tree.)
- If no back edge is found, then there is no directed cycle.

Example

In the DFS execution above, e_8 is a back edge, and it closes the cycle (e_6, e_7, e_8).

Cycle Detection





- A correct approach: Perform directed DFS.
- If a *back edge* is found, then it closes a cycle. (A back edge is an edge whose destination is an ancestor of its origin in the spanning tree.)
- If no back edge is found, then there is no directed cycle.

Example

In the DFS execution above, e_8 is a back edge, and it closes the cycle (e_6, e_7, e_8) .

Strong Connectivity



Definition

A directed graph $G = (V, E)$ is strongly connected if u reaches v and v reaches u for every $u, v \in V$.


Examples

G is *not* strongly connected because v_1 is not reachable from v_2 .
 H is strongly connected.


Strong Connectivity

- How can we check whether a graph $G(V, E)$ is strongly connected?
- Check whether $\text{DFS}(v)$ reaches all the vertices in V for each $v \in V$.
- Running time: $O(n(m + n))$.
- A better approach: (see next slide)

Strong Connectivity



G



G'

Strong Connectivity Algorithm

- Choose a vertex $s \in V$
 - If $\text{DFS}(G, s)$ does not reach all the nodes, return FALSE.
 - Let G' be the graph obtained by reversing all the edges in G .
 - If $\text{DFS}(G', s)$ does not reach all the nodes, return FALSE.
 - Otherwise return TRUE.
-
- This algorithm runs in time $O(m + n)$.



Directed Acyclic Graphs

Definition

A *directed acyclic graph (DAG)* is a directed graph with no directed cycle.

Definition

A *topological ordering* of a directed graph $G = (V, E)$ is an ordering $\{v_1, \dots, v_n\}$ of its vertices such that $i < j$ whenever $(v_i, v_j) \in E$.



Two topological orderings of the same DAG.


Directed Acyclic Graphs

Proposition

A directed graph G has a topological ordering if and only if it is acyclic.


- We need to prove two statements:
 - \Leftarrow If G has a topological ordering, then it is acyclic.
 - \Rightarrow If G is acyclic, then it has a topological ordering.
- Proof of \Leftarrow : done in class.
- Proof of \Rightarrow : We make a constructive proof by showing how to construct a topological ordering of G .

Directed Acyclic Graphs




- G must have at least one vertex of indegree 0 because:
- If it were not the case, we could trace a path backwards from a starting vertex, and we would visit the same vertex twice. So G would have a directed cycle.
- We let v_1 be this vertex, and we delete this vertex from G .

Directed Acyclic Graphs




- We remove v_1 from G , and the resulting graph must still be acyclic.




- We find a vertex v_2 of indegree 0 in this new graph, and repeat the process.

Directed Acyclic Graphs




- Now there are two choices for v_4 . We choose the bottom vertex.



- We find a vertex v_2 of indegree 0 in this new graph, and repeat the process.

Directed Acyclic Graphs



- In the end, we find a topological ordering of G .
- We just showed how to find a topological ordering for any graph G .
- It also gives us an algorithm to construct it.
- We call this algorithm *topological sorting*.

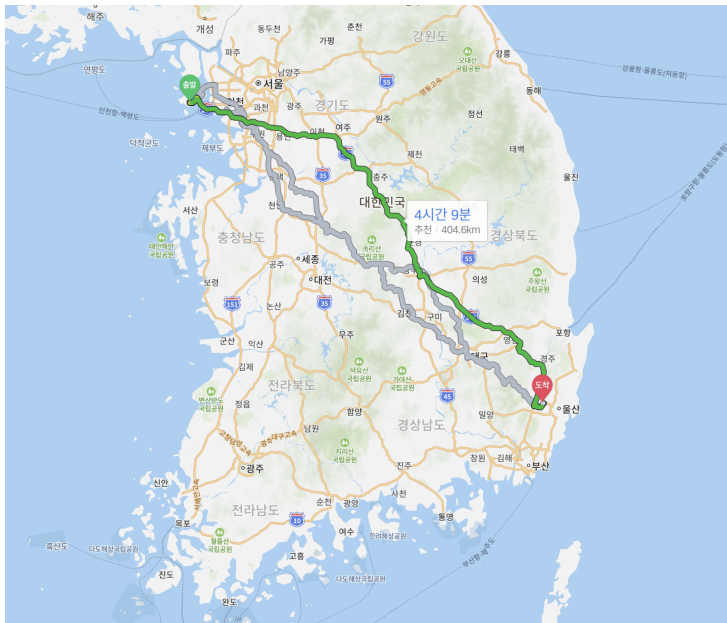
Topological Sorting

```
procedure TOPOLOGICALSORT( $G$ )  
   $S \leftarrow$  an empty stack  
  for all  $u$  in  $G.vertices()$  do  
     $incounter(u) \leftarrow indeg(u)$   
    if  $incounter(u)=0$  then  
       $S.push(u)$   
  
   $i \leftarrow 1$   
  while  $S \neq \emptyset$  do  
     $u \leftarrow S.pop()$   
     $v_i \leftarrow u$   
     $i \leftarrow i + 1$   
    for all outgoing edges  $(u, w)$  of  $u$  do  
       $incounter(w) \leftarrow incounter(w)-1$   
      if  $incounter(w)=0$  then  
         $S.push(w)$ 
```


Topological Sorting

- We can initialize the `incounter()` variables in $O(n + m)$ time overall by performing a traversal of G .
- Other than that, we make $O(m + n)$ stack operations.
- So `TOPOLOGICALSORT` runs in $O(m + n)$ time.
- Remark: This algorithm also allows us to decide whether G is acyclic. Indeed, if it fails to sort all the vertices, then the graph must have a cycle.
- Application: A number of tasks have to be performed.
- We represent them as a graph, with an edge (i, j) if task i needs to be completed before task j .
- Then a topological ordering gives you a possible order in which the tasks can be performed.

Shortest Paths




Weighted Graphs



- A *weighted graph* is a graph that has a numeric (for example, integer) label $w(e)$ associated with each edge e , called the *weight* of edge e .

Weighted Graphs




- The *length* of a path is the sum of the lengths of the edges of the path.

Example

The path shown in the graph above has length 13.

Shortest Paths




- The *distance* from a vertex v to a vertex u in G , denoted $d(v, u)$, is the length of a minimum length path (also called *shortest path*) from v to u , if such a path exists.

Example

In the graph above, $d(v_1, v_4) = 9$.


Dijkstra's Algorithm



$Q = (v_1, v_2, v_3, v_4, v_5)$

priority queue with d -values as keys


Dijkstra's Algorithm



$$Q = (v_2, v_3, v_4, v_5)$$

remove the minimum v_1 from Q

Dijkstra's Algorithm




$$Q = (v_2, v_3, v_4, v_5)$$

relax the edge (v_1, v_2) :

$$d_2 \leftarrow \min(d_2, d_1 + 10)$$

Dijkstra's Algorithm




$$Q = (v_2, v_3, v_4, v_5)$$

relax the edge (v_1, v_3) :

$$d_3 \leftarrow \min(d_3, d_1 + 5)$$

Dijkstra's Algorithm




$Q = (v_3, v_2, v_4, v_5)$

relax the edge (v_1, v_3) :

$$d_3 \leftarrow \min(d_3, d_1 + 5)$$


Dijkstra's Algorithm



$$Q = (v_2, v_4, v_5)$$

remove the minimum v_3 from Q


Dijkstra's Algorithm



$$Q = (v_2, v_4, v_5)$$

relax edge (v_2, v_3)


Dijkstra's Algorithm



$Q = (v_2, v_4, v_5)$

relax edge (v_3, v_4)


Dijkstra's Algorithm



$$Q = (v_5, v_2, v_4)$$

relax edge (v_3, v_5)


Dijkstra's Algorithm



$$Q = (v_2, v_4)$$

remove minimum v_5 from Q


Dijkstra's Algorithm



$Q = (v_2, v_4)$

relax edge (v_5, v_4)


Dijkstra's Algorithm



$$Q = (v_4)$$

remove minimum v_2 from Q


Dijkstra's Algorithm



$Q = (v_4)$

relax edge (v_2, v_4)


Dijkstra's Algorithm



$Q = ()$

remove minimum v_4 from Q

Dijkstra's Algorithm



$$Q = ()$$

All the distances from v_1 have been computed.

Dijkstra's Algorithm

- *Dijkstra's algorithm* solves the *single-source* shortest path problem: It computes the distance from a vertex s to all the other vertices in G .
- It requires the weights to be non-negative. It works for both directed and undirected graphs.
- For each vertex u , it maintains the length $D[u]$ of the shortest path to u that we have found *so far*. Initially, $D[s] = 0$ and $D[u] = \infty$ for all $u \neq s$.
- We maintain a set C of vertices (called the *cloud*) which is initially empty. At each iteration of the main loop, we insert the vertex $u \notin C$ with the smallest label $D[u]$ into C .

Dijkstra's Algorithm

- Then we update the labels of the vertices z adjacent to u and that are not in C using an *edge relaxation* operation:

Relaxation of edge (u, z)

if $D[u] + w((u, z)) < D[z]$ **then**
 $D[z] \leftarrow D[u] + w((u, z))$

- The vertices that are not in the cloud C are maintained in a priority queue Q , with the labels $D[u]$ as keys.

Dijkstra's Algorithm

Pseudocode

procedure SHORTESTPATH(G, s)

 set $D[s] \leftarrow 0$ and $D[u] \leftarrow \infty$ for all $u \neq s$

 let Q be a priority queue recording all the vertices u with keys $D[u]$

while $Q \neq \emptyset$ **do**

$u \leftarrow Q.\text{removeMin}()$

for each edge (u, z) such that $z \notin Q$ **do**

if $D[u] + w((u, z)) \leq D[z]$ **then**

$D[z] \leftarrow D[u] + w((u, z))$

 change to $D[z]$ the key of z in Q

return the label $D[u]$ of each vertex u

Analysis

- Each vertex is removed exactly once from Q , so there are n iterations of the **while** loop.
- So we make n `removeMin()` operations, which takes $O(n \log n)$ time overall if Q is implemented by a heap.
- Each edge is relaxed at most once, so there are m iterations of the **for** loop.
- Each iteration of the for loop takes $O(\log n)$ time for updating the key of z .

Proposition

Dijkstra's algorithm solves the single-source shortest path problem for a graph with non-negative weights in time $O((m + n) \log n)$.

Remarks

- At each iteration, Dijkstra's algorithm picks the node that is closest to s .
- This is an example of a *greedy algorithm*.
- More generally, greedy algorithms repeatedly select the best choice available at each iteration.
- The proof of correctness is non-trivial. I do not give it due to lack of time. See CSE331: Introduction to Algorithms.
- The pseudocode above does not show how to compute a shortest path: It only computes the distance from s to u .
- In order to find a shortest path from s to v , we can simply remember after each relaxation operation the node u from which the shortest path to z came. (See next slide.)
- Then we can recover the path by following the references $\text{prev}[z]$ backwards from z .

Dijkstra's Algorithm

Pseudocode

procedure SHORTESTPATH(G, s)

set $D[s] \leftarrow 0$ and $D[u] \leftarrow \infty$ for all $u \in V \setminus \{s\}$

set $\text{prev}[u] \leftarrow \text{NULL}$ for all $u \in V$

let Q be a priority queue recording all the vertices u with keys $D[u]$

while $Q \neq \emptyset$ **do**

$u \leftarrow Q.\text{removeMin}()$

for each edge (u, z) such that $z \notin Q$ **do**

if $D[u] + w((u, z)) \leq D[z]$ **then**

$D[z] \leftarrow D[u] + w((u, z))$

 change to $D[z]$ the key of z in Q

$\text{prev}[z] \leftarrow u$

return $D[u]$ and $\text{prev}[u]$ for each vertex u